

Pavel Málek



Score a (row) goal and beat a query optimizer

Sponsors



Bio



- 10+ years of experience with SQL Server
- Czech PASS leader
- SQL Saturday Prague co-organizer



- Lead SQL Server DBA in largest Czech bank for 6+ years



- Currently newbie at Joyful Craftsmen

 [@malekpav](https://twitter.com/malekpav)

 [LinkedIn](https://www.linkedin.com/in/malekpav)

 malekpav@gmail.com




Before we begin...



Challenge for you!

Try to write a query which will show you 10 last trades (*trd.TradesArch*) for each market trader (*trd.Parties*)?



It will be used on primary site for traders => it has to be faaaast!





Solution

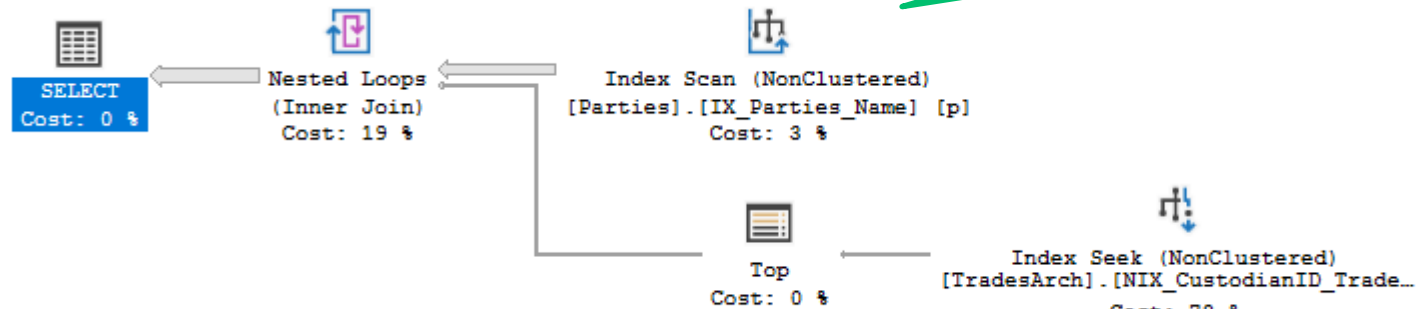
Nature of SQL

- SQL is a declarative language!



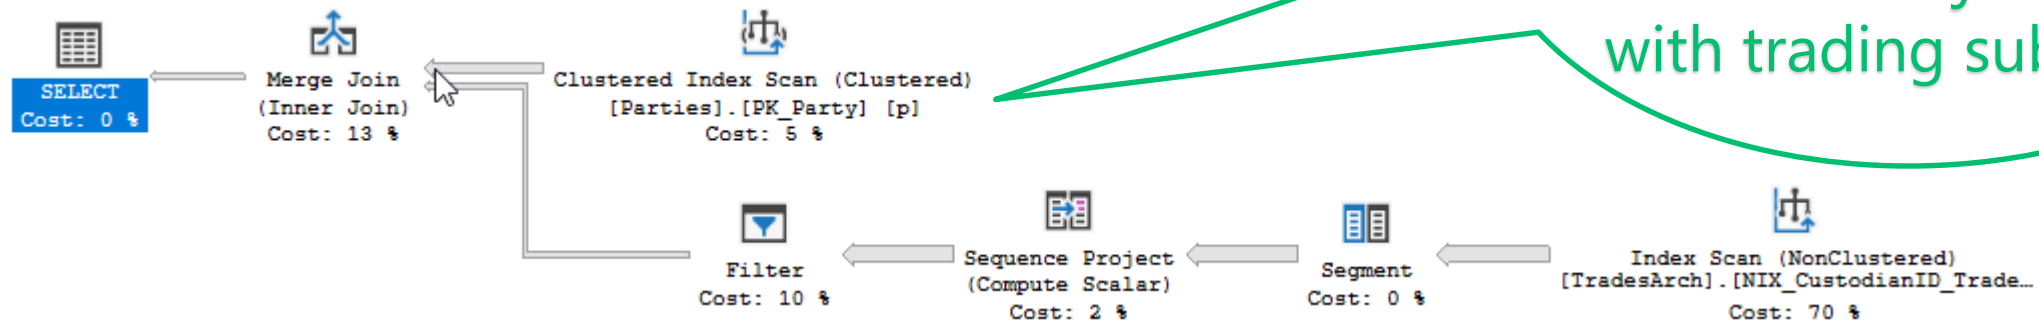
But how much “declarative”?

Cross apply



For every subject trading on market find the most recent related trade

RowNb



Go through all the trades and show only those who are the most recent joined with trading subjects






Query Optimizer is in charge, but...

- Sometimes QO simply cannot have correct inputs

SQL Server Customer Advisory Team

When To Break Down Complex Queries

September 9, 2013 by [AzureCAT](#) // [6 Comments](#)


 Share 1  0  4

★★★★★

- OR logic in the WHERE clause
- Aggregations in intermediate results sets
- A large number of very complex joins
- A CASE clause in the WHERE or JOIN clause





Query optimizer having troubles

RE: Long responses in BO 



Pavel Málek <pmalek@csas.cz>

15.11.17 

 Odpovědět všem 

komu: Jan, sITS, Martin 

There are currently only two long running queries:

```
SELECT TOP (80) * FROM trd.TradesView ORDER BY [Id] DESC, [Version] DESC
```

```
SELECT TOP (80) * FROM trd.TradesArchView ORDER BY [Id] DESC, [Version] DESC
```

Query plan is attached. Very first impressions - nested loop everywhere, estimated nr of rows is not accurate. Wrong stats? I am going to update them...

P.

```
SELECT TOP (80) * FROM trd.TradesView ORDER BY [Id] DESC, [Version] DESC
```


```
SELECT TOP (80) * FROM trd.TradesArchView ORDER BY [Id] DESC, [Version] DESC
```






Demo


Histogram aligning



(a) Original histograms.



(b) Step 1: Align buckets.



(c) Step 2: Process each bucket.

Figure 2.8: Cardinality estimation of join predicates using histograms.

Nicolas's Bruno ([Software Developer, Microsoft](#)) work "*Statistics on Query Expressions in Relational Database Management Systems*" COLUMBIA UNIVERSITY, 2003.

"Mere mortals" reading:
Hugo Kornelis – SQL Sat [presentation](#)
Dimo – [queryprocessor.com](#)



Other possibilities?

Simple Join: Simple join ignores histograms from the two relations, and employs a containment model that simply assumes all distinct values from one child are contained in the other (originally introduced in [8]). This simplistic strategy typically sacrifices estimation quality for faster optimization time, providing a balance that better suits some workloads.



Demo Takeaways

- FKs are good
- Indexes on FKs is almost always good
- But don't expect them to be "always" used for row estimation!
- **Updates of stats via full scan matters!!!**
- TF 2301 sometimes helps
- Carefully with TF 2371

[Sign in](#) to post a comment.

Posted by **Paul White NZ** on 6/3/2013 at 11:48 PM

Not common enough? Many production databases use foreign keys (yes, really) and better cardinality estimation for joins across such a relationship would be of wide benefit. If you have plans to address the underlying issue in other ways, please say so. Otherwise, "Won't Fix" seems like a very curious decision indeed.



Demo case continues

What if I don't want to use any traceflags

TOP operator introduced an interesting behavior

What the hell "a row goal" is?



Case Trading System#2: Key Takeaways

Via TOP + ORDER BY you are able to create your own "blocks"

Via TOP (@x) OPTIMIZE FOR you are able to decrease estimated nb of rows

Via MANY() you are able to enlarge estimated nb of rows

CREDITS: Adam Machanic



Another example of fixed QO behavior



Summary

Now you know:

How much statistics matter

How joins are calculated

How to force QO to do what you want to

(TOP blocks, MANY(), TOP (@x) + OPTIMIZE FOR)

CROSS APPLY vs ROWNB()

What are the possibilities for MAX() and MIN()



Thank you!



Evaluations

Event Evaluation

<http://www.sqlsaturday.com/764/EventEval.aspx>

Session Evaluation

<http://www.sqlsaturday.com/764/Sessions/SessionEvaluation.aspx>

